

Distributed Systems

Unit-9: Security

Compiled by Prashant Gautam

Why Security in Distributed Systems?

- Distributed Systems rely on sharing of resources across different networked entities
- Most vital/secret data handled by distributed components
- A single security flaw compromises the whole system
 - Malware and viruses can spread from one part of the system to another easily
 - Users across the world may have access to the system
 - Cyber criminals, hackers
- Security lapses result in
 - Loss of confidence, Claims for damages, Loss of privacy

Overview

Introduction

- Threats, policies and mechanisms
- Cryptography

Secure Channels

- Authentication
- Message Integrity and Confidentiality

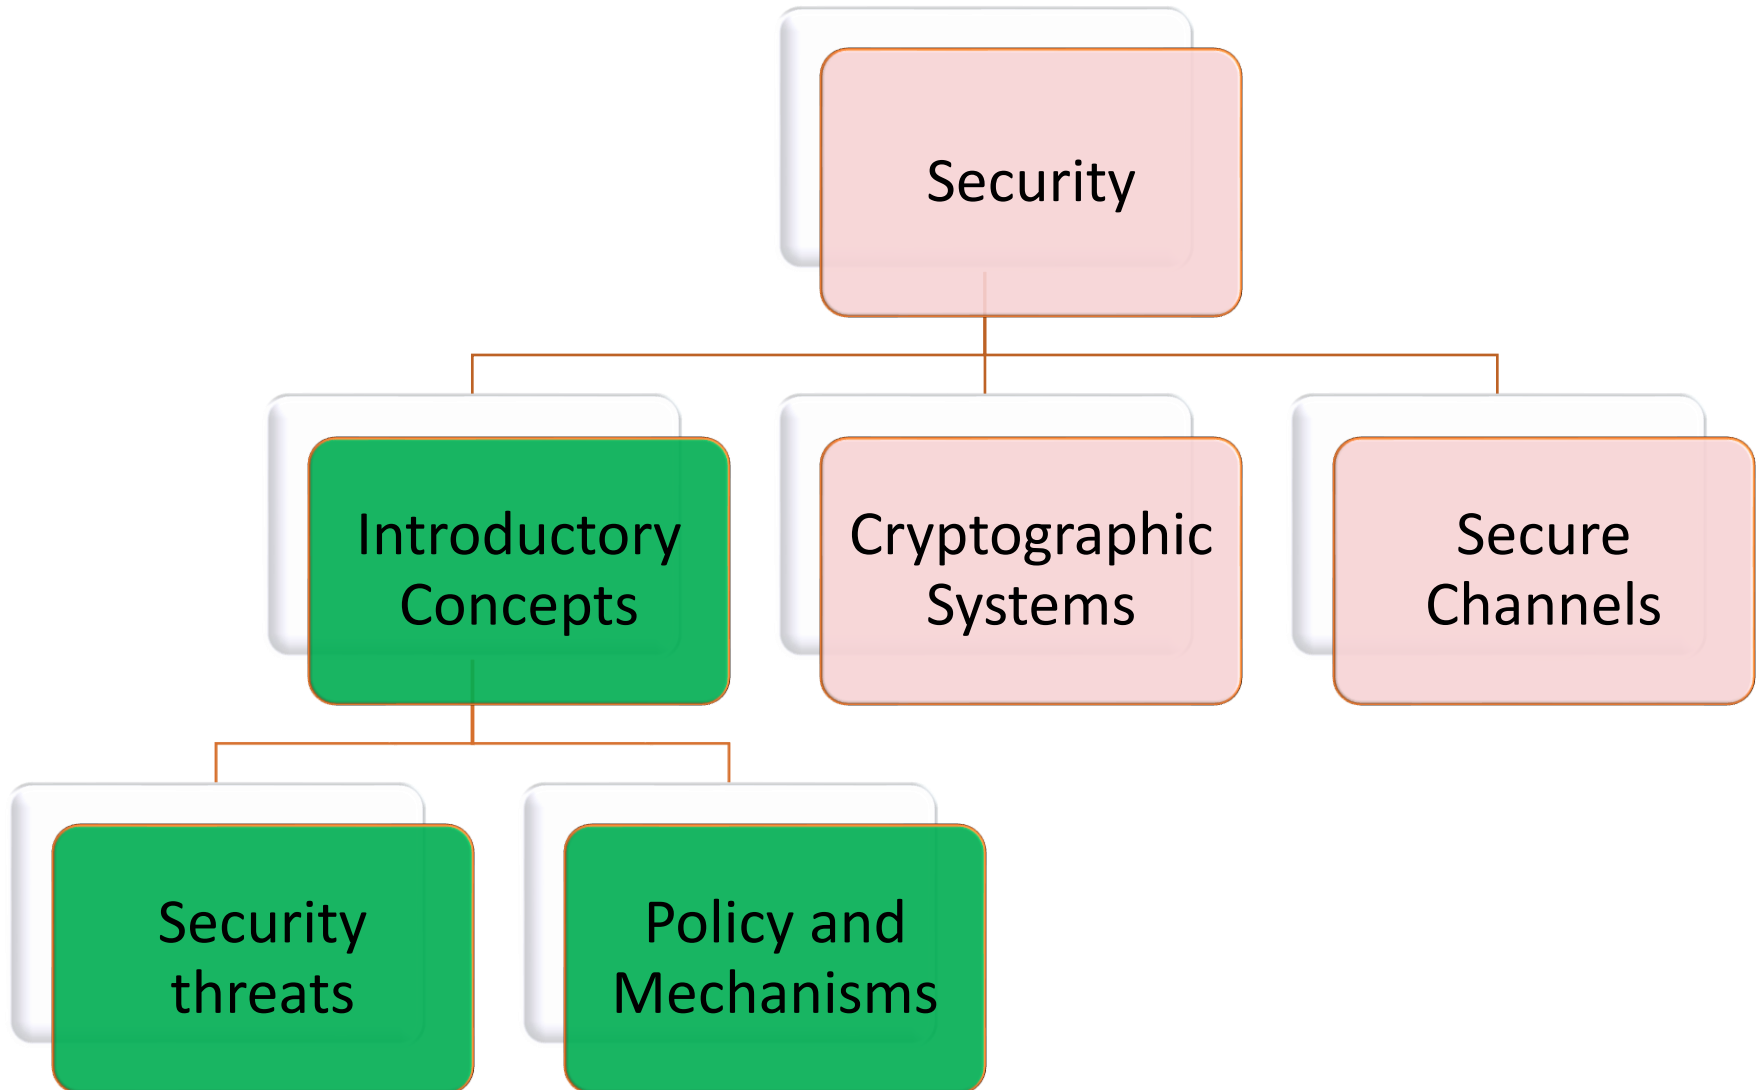
Access Control

- Access Control Matrix
- Protection Domains

Security Management

- Key Management
- Authorization Management

Overview



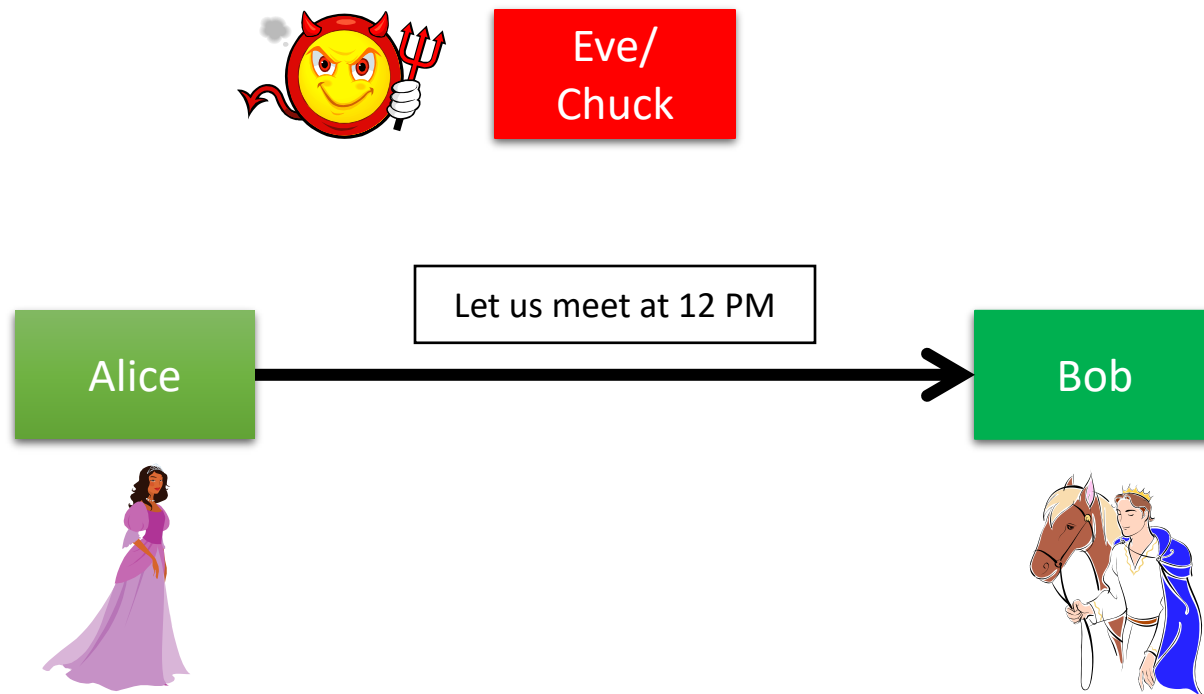
Introduction to Security

- What services do you expect from secure Distributed Systems?
- Secure DS should provide
 - Confidentiality of Information
 - Information is disclosed only to authorized parties
 - Integrity of Information
 - Alterations to system's assets is made only in an authorized way

Secure DS are immune against possible security threats that compromise confidentiality and integrity

Security Threats


- What are the security threats when two entities communicate?



Types of Security Threats (1)


1. Interception

- Unauthorized party has gained access to a service or data
- Example: Illegal copying of files, Eavesdropping over network



2. Interruption

- Services or data become unavailable, unusable or destroyed
- Example: Denial-of-Service (DoS) Attacks





Types of Security Threats (2)

3. Modification

- Unauthorized changing of data or tampering with services
- Example: Changing a program to secretly log the activities

4. Fabrication

- Additional data or activity is generated that would normally not exist
- Example: Replay attacks



Security Policy and Mechanisms

- To build a secure DS, we need to

Formulate **Security Policies**

- Specifies what is to be done

Describe the security requirements:

- which actions the entities are allowed to take
- which actions are prohibited

Build **Security Mechanisms**

- Specifies how policies are implemented

Implement mechanisms to:

- Protect data transferred
- Verify identity of an entity
- Secure access permissions

Security Mechanisms

Four core components of security mechanisms

Cryptographic Algorithms and Secure Channels

1. Encryption

- Transform the data to something that attacker cannot understand

2. Authentication

- Verifies the claimed identity of the user, host or other entity

Access Control

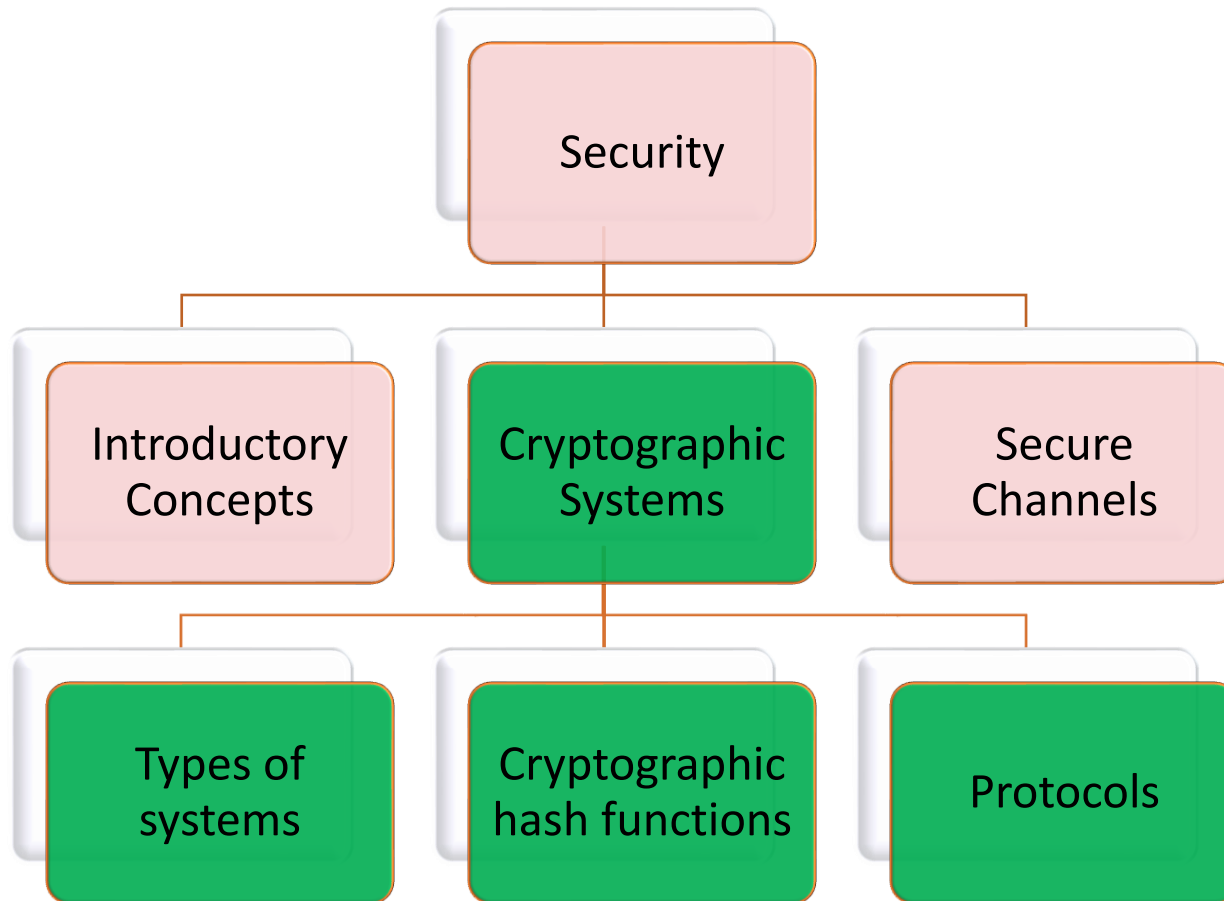
3. Authorization

- Verifies if the entity is authorized to perform an operation

4. Auditing

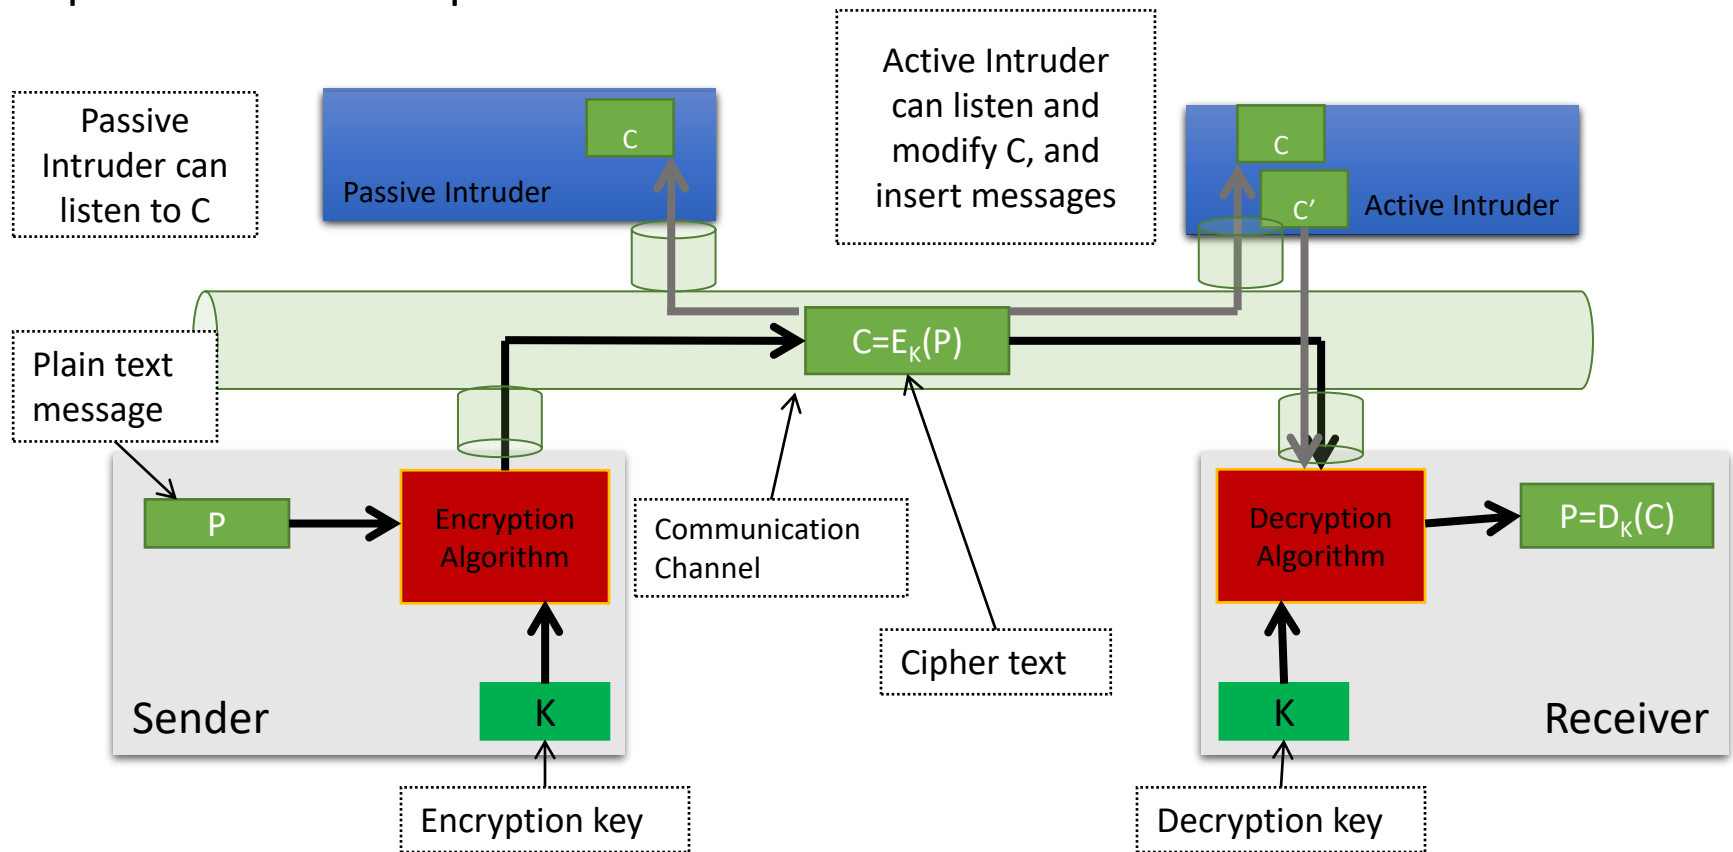
- To trace which clients accessed what, and which way

Overview



Cryptographic systems

- Cryptography is the study of techniques for secure communication in the presence of third parties



$C = E_K(P) \rightarrow C$ is obtained by encrypting the plain text P with key K

$P = D_K(C) \rightarrow P$ is obtained by decrypting the cipher text C with key K

Types of Cryptographic Systems


- Two types
 1. Symmetric Cryptosystem (Shared-key system)
 2. Asymmetric Cryptosystem (Public-key system)

Symmetric Cryptographic System

- Same key is used to encrypt and decrypt the data

$$P = D_K(E_K(P))$$

- The shared key $K_{A,B}$ between Alice A and Bob B should be kept secret
- Also known as *secret-key* or *shared-key* systems



Public-key Cryptographic System


- The key for encryption (K_E) and decryption(K_D) are different
- But, K_E and K_D form a unique-pair

$$P = D_{K_D}(E_{K_E}(P))$$

- One of keys is made public, and another made private
- Denotation: Public key of A = K_A^+ ; Private key of A = K_A^-
- Public-key systems can be used for:
 - Encryption
 - Authentication

Encryption in Public-key system


- Scenario: Alice (A) wants to send to Bob (B)
- Problem: Only Bob should be able to decrypt the message
- Approach:
 - At A: Encrypt using B's public key
 - At B: Decrypt using B's private key



Drawback: How does 'B' know that 'A' sent the message?

Authentication in Public-key system


- Scenario: Alice (A) wants to send a message to Bob (B)
- Problem: Bob wants to make sure that message came from Alice (and not from some intruder)
- Approach:
 - At A: Encrypt using A's private key
 - At B: Decrypt using A's public key



Drawback: How to ensure that ONLY 'B' gets the message?

Combining encryption and authorization

- Scenario: Alice (A) wants to send a message to Bob (B)
- Many algorithms use a combination of the above two methods to:
 - Ensure that only Bob can decrypt the message
 - Bob can verify that Alice has sent the messages



- Approach: Encrypt/Decrypt using a combination of keys
 - A widely used method in many secure algorithms

• What happens if intruder 'C' modifies message 'm' sent by 'A'?

Some part of message 'm' should contain data that verifies the message

Cryptographic Hash Functions

- A hash function $H(m)$ maps an input message m to a hash value h
 - Message m is of any arbitrary length
 - Hash h is fixed length

$$h = H(m)$$

- Often, h is called as the “*message digest*” of m
- How does a cryptographic hash function differ from a regular hash function?

Properties of Cryptographic Hash Functions

1. One-Way Function

- Finding hash h from m is easy, but not vice-versa
- Computationally infeasible to find m that corresponds to a given hash h

2. Weak-collision resistance

- Given a message, it is hard to find another message that has the same hash value
- Given m and $h = H(m)$, it is computationally infeasible to find $m' \neq m$ such that $H(m) = H(m')$

3. Strong-collision resistance

- Given a hash function, it is hard to find two messages with the same hash value
- Given $H(.)$, it is computationally infeasible to find two messages m and m' such that $H(m) = H(m')$

Encryption/Decryption Functions

- Recall: An encryption function $E_K(m_p)$ encrypts a plain-text message m_p to an cipher-text message m_c using a key K

$$m_c = E_K(m_p)$$

- Similarly, decryption function $D_K(m_c)$:

$$m_p = D_K(m_c)$$

- Encryption/Decryption Functions have the same properties as Cryptographic Hash Functions
 - One-way functions
 - Weak and Strong collision resistance

Additional Properties of Encryption/Decryption Functions

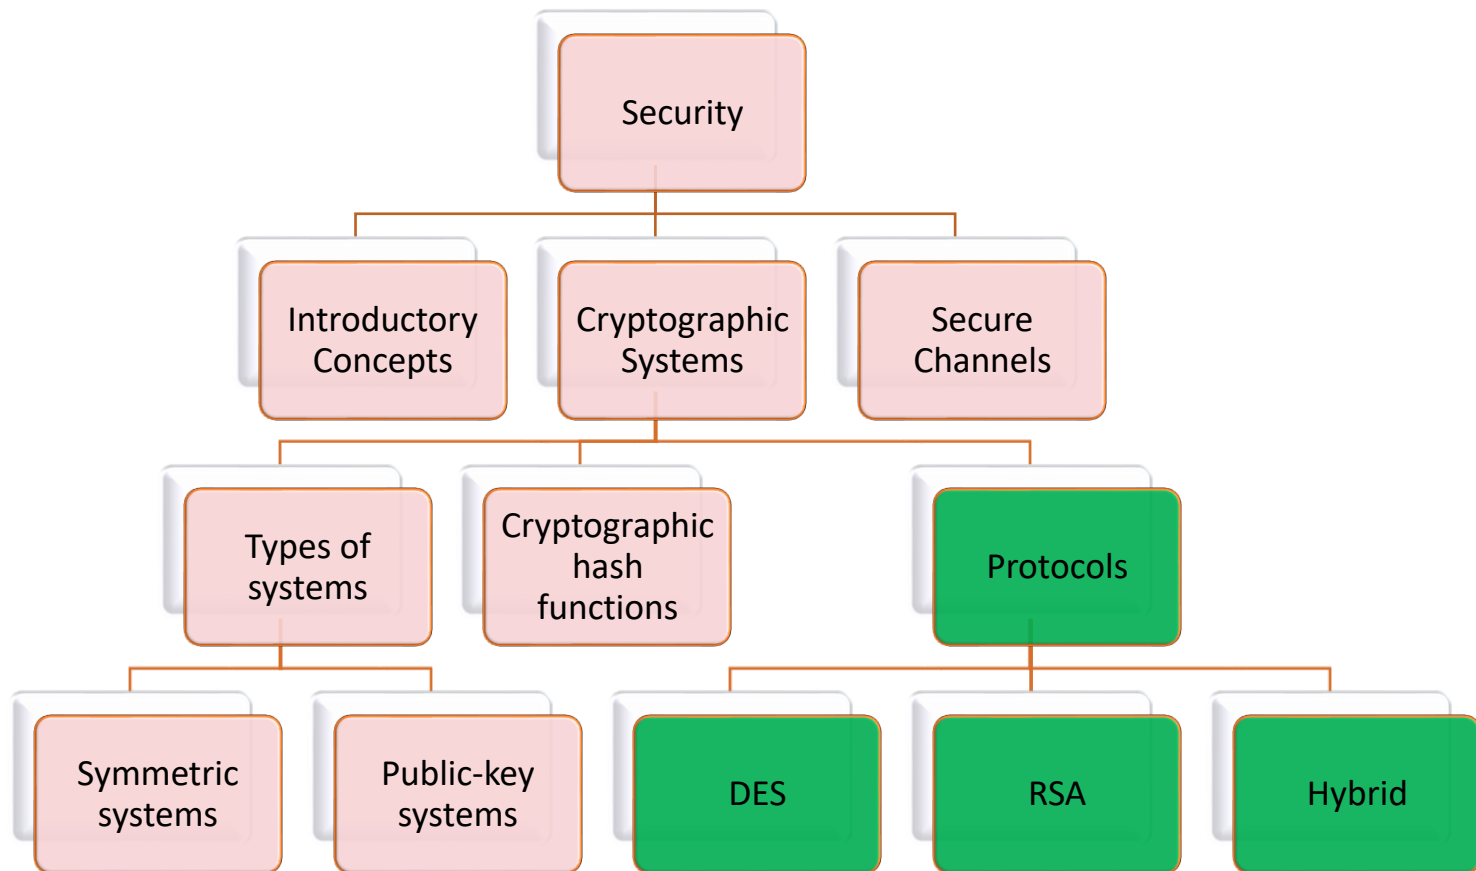
- Infeasible Key Extraction *

- Given a plain-text m_p and its cipher-text m_c , it is hard to find the key K such that $m_c = E_K(m_p)$

- Key Collision Resistance *

- Given a plain-text and a key K , it is hard to find another key K' that produces the same cipher-text
- Given a plain-text m_p and a key K such that cipher-text $m_c = E_K(m_p)$, it is hard to find another key K' that produces the same cipher-text m_c

Overview


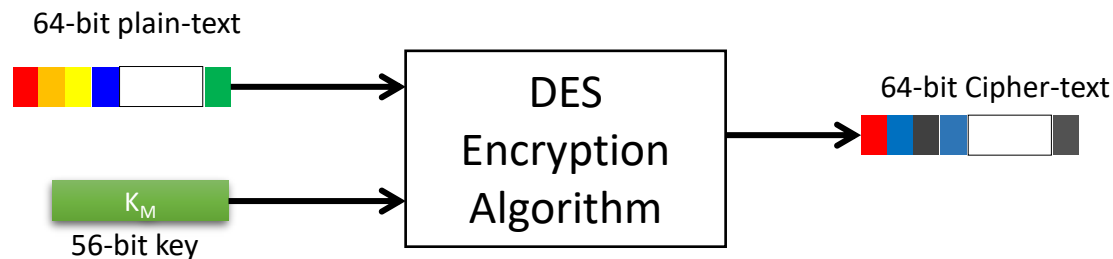


Protocols in Cryptosystems

- We will study three protocols in cryptosystems
 - Data Encryption Standard (DES)
 - Encryption/Decryption in Symmetric Cryptosystems
 - RSA protocol
 - Encryption/Decryption in Public-Key cryptosystems
 - Hybrid Cryptographic protocol
 - A combination of Symmetric and Public-Key based system

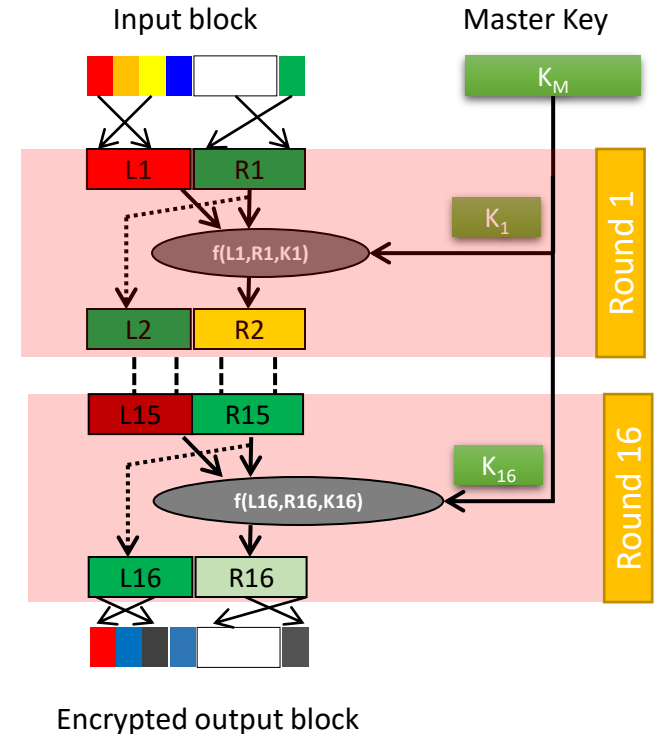
DES Protocol

- Purpose: Encryption/Decryption in Symmetric Cryptosystems
- Encryption and Decryption relies on a 56-bit master key (K_M)
- Operates on a 64-bit block of input data to encrypt/decrypt



DES Encryption Algorithm

1. Permute a 64-bit block
2. 16 rounds of identical operations
3. In each round i
 - i. Divide the block into 2 halves L_i and R_i
 - ii. Extract 48-bit key K_i from K_M
 - iii. Mangle the bits in L_i and R_i using K_i to produce R_{i+1}
 - iv. Extract R_i as L_{i+1}
4. Perform an inverse permutation on the block L_{16} - R_{16} to produce the encrypted output block



Discussion about DES

- DES encryption and decryption is relatively fast
- DES has disadvantages of a Symmetric Cryptosystem
 - Requires sender and receiver to exchange K_M
 - For N-user system, DES needs $N(N-1)/2$ master key pairs
- History of DES:
 - DES was invented in 1974
 - In 1997, it was shown that DES can be easily cracked using brute-force attacks
- Triple-DES is in use in some systems
 - It applies DES three times using two keys

RSA protocol


- Invented by Rivest, Shamir and Adleman (RSA) as a protocol for Public-key systems
- Approach:
 1. Choose two very large prime numbers, p and q
 2. Compute $n = pq$
 3. Compute $z = (p-1)(q-1)$
 4. Choose a number e that is relatively prime to z
 - e is co-prime to z
 5. Compute the number d such that $de \% z = 1$
 - This is equivalent to finding $de = 1 + kz$ for some integer k
- Depending on the requirement, d and e can be used as public and private keys
 - d is used for decryption; e is used for encryption

Example*
$p=7; q=19$
$n = 7*19 = 133$
$z = 6*18 = 108$
$e=5$
$d=65$ for $m=325$

Example: RSA protocol for encryption (1)

- Scenario: Alice (A) wants to send to Bob (B)
- Problem: Only Bob should be able to decrypt the message
- Given d and e are the two keys computed by RSA, which row indicates correct choice of keys?

Correct/ Incorrect	Alice		Bob	
	Private K_A^-	Public K_A^+	Private K_B^-	Public K_B^+
X	d	e		
X	e	d		
Correct			d	e
X			e	d



Example: RSA protocol for encryption (2)

At the sender:

- Split the message into fixed-length blocks of size s ($0 \leq s < n$)
- For each block m_i
 - Sender calculates the encrypted message c_i such that $c_i = m_i^e \pmod{n}$
 - Send c_i to the receiver

At the receiver:

- Receive c_i from sender
- For each block c_i
 - Compute actual message $m_i = c_i^d \pmod{n}$
- Merge all c_i 's to obtain the complete message

Calculated Values

$p=7$; $q=19$; $n=133$; $d=65$; $e=5$

Example
$s = 132$
$m_i = 6$
$c_i = 6^5 \pmod{133}$ $= 62$
$m_i = 62^{65} \pmod{133}$ $= 6$

Discussion about RSA

- RSA has advantages of Public-key system
 - Harder to break the code
 - For a N-user system, RSA needs only $2N$ keys
- Computation time of RSA is much larger than DES
 - Approximately 100-1000 times slower

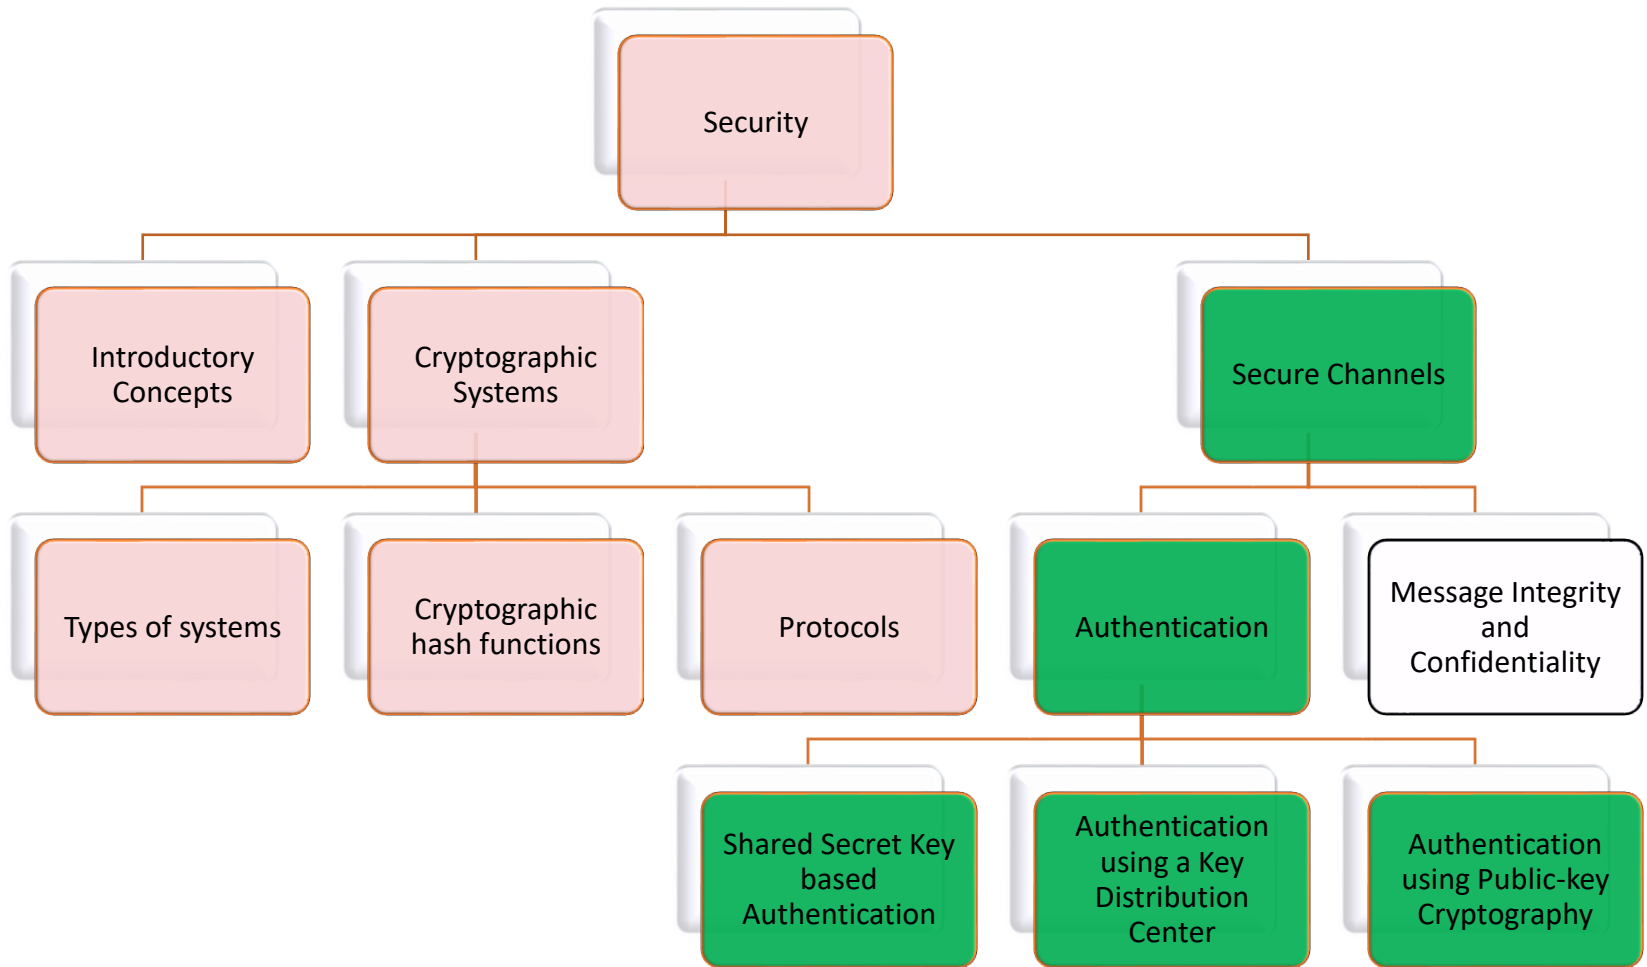
Hybrid Cryptographic protocols

- Large scale distributed systems use a combination of symmetric and public-key protocols
- Leveraging the advantages of both schemes
 - Encryption based on Public-key are more secure
 - Authenticate using RSA
 - Exchange the “secret key” using RSA for a session
 - Encryption based on Symmetric keys are faster
 - Exchange large data using the above “secret key”

Beyond Cryptographic Mechanisms

- Many users, clients and servers need to dynamically send messages in a distributed system
 - How can an end-to-end secure distributed system be built using the cryptographic mechanisms?
 - How can each message and user be protected against security threats?
 - How do clients and processes authenticate?
 - What protocols are needed for these? What is their complexity?

Overview



Secure Channels

- A *Secure Channel* is an abstraction of secure communication between communication parties in a DS
- A Secure Channel protects senders and receivers against:
 - Interception
 - By ensuring *confidentiality* of the sender and receiver
 - Modification and Fabrication of messages
 - By providing *mutual authentication* and *message integrity* protocols
- We will study
 - Authentication
 - Confidentiality and Message Integrity

Authentication

- Consider a scenario where Alice wants to set up a secure channel with Bob
 - Alice sends a message to Bob (or trusted third party) for mutual authentication
 - Message integrity should be ensured for all communication between Alice and Bob
 - Generate a “session key” to be used between Alice and Bob
 - Session-keys ensure integrity and confidentiality
 - When the channel is closed, the session key is destroyed

Types of Mutual Authentication Protocols

1. Shared Secret Key based Authentication
2. Authentication using a Key Distribution Center
3. Authentication using Public-key Cryptography

Types of Mutual Authentication Protocols

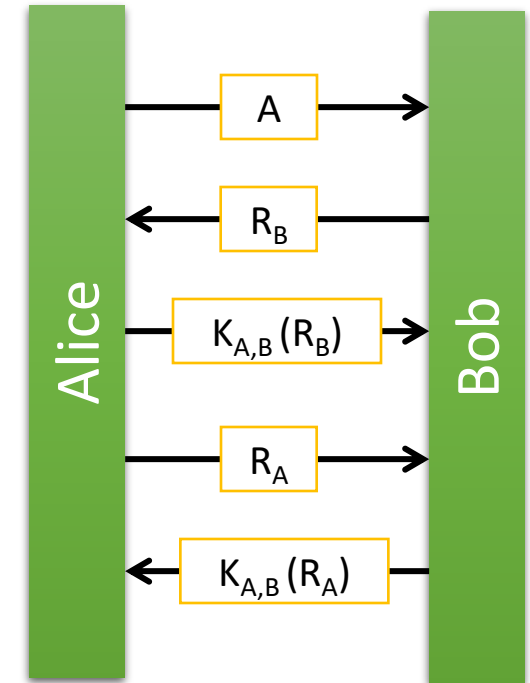
1. Shared Secret Key based Authentication
2. Authentication using a Key Distribution Center
3. Authentication using Public-key Cryptography

Shared Secret Key based Authentication

- The scheme is also known as “*Challenge-Response protocol*”
- Let $K_{A,B}$ be the shared secret key between Alice and Bob

The Challenge-Response Protocol

1. ‘A’ sends her identity to ‘B’
2. ‘B’ sends a challenge R_B back to ‘A’
3. ‘A’ responds to the challenge by encrypting R_B with $K_{A,B}$ (denoted by $K_{A,B}(R_B)$), and sending it back to ‘B’
4. ‘A’ challenges ‘B’ by sending R_A
5. ‘B’ responds to the challenge by sending the encrypted message $K_{A,B}(R_A)$

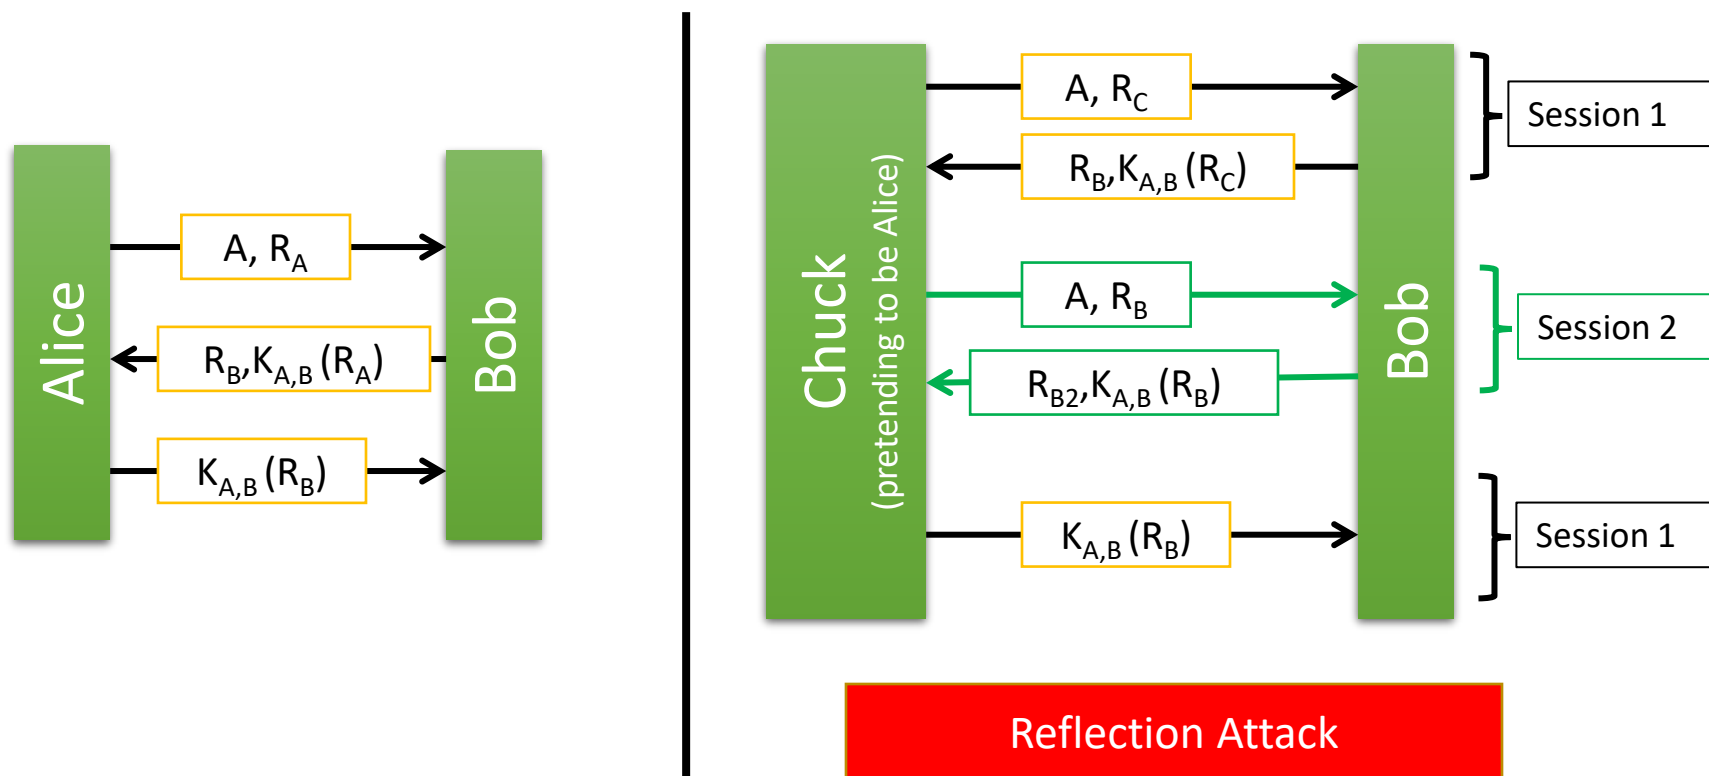


A and B are mutually authenticated

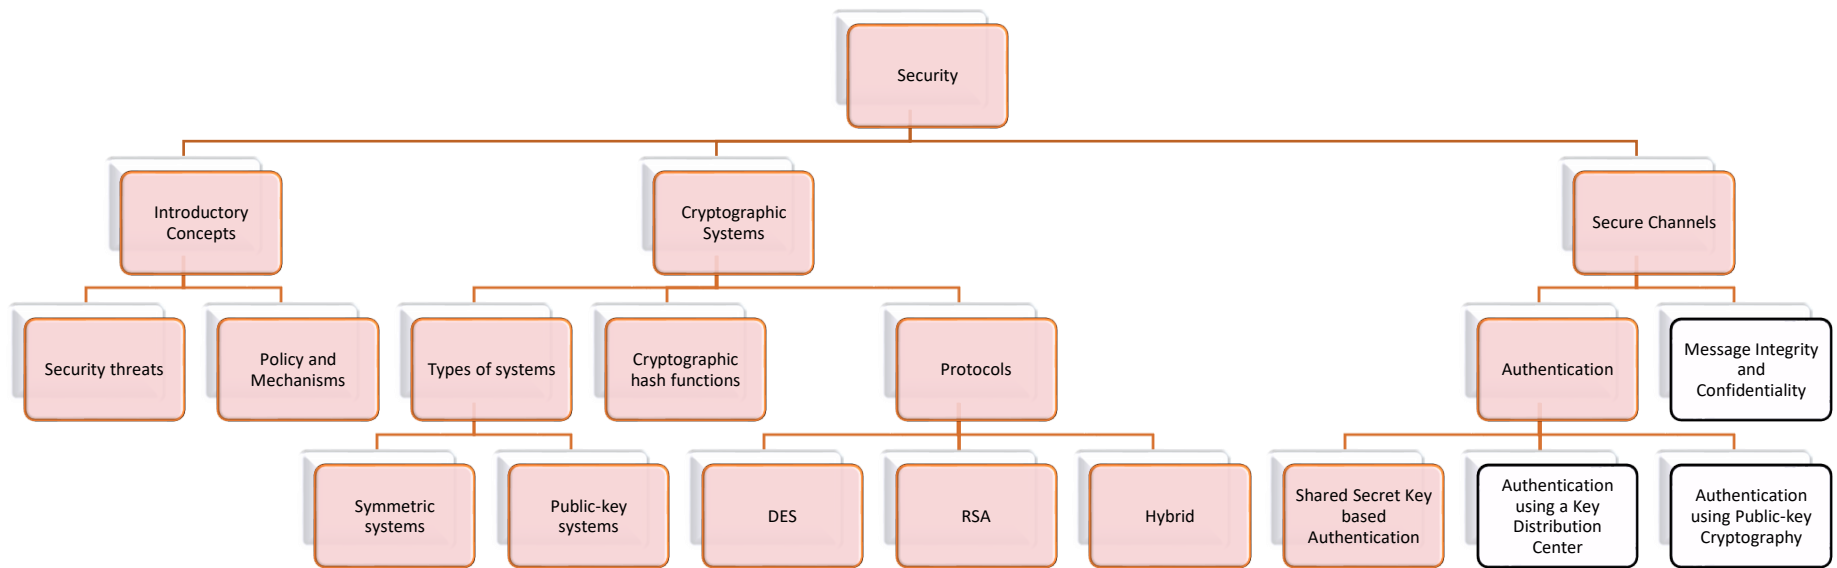
(?)

A Possible Optimization

- Will the below 3-step protocol work?



Summary



Overview

Introduction

- Threats, policies and mechanisms
- Cryptography

Secure Channels

- Authentication
- Message Integrity and Confidentiality

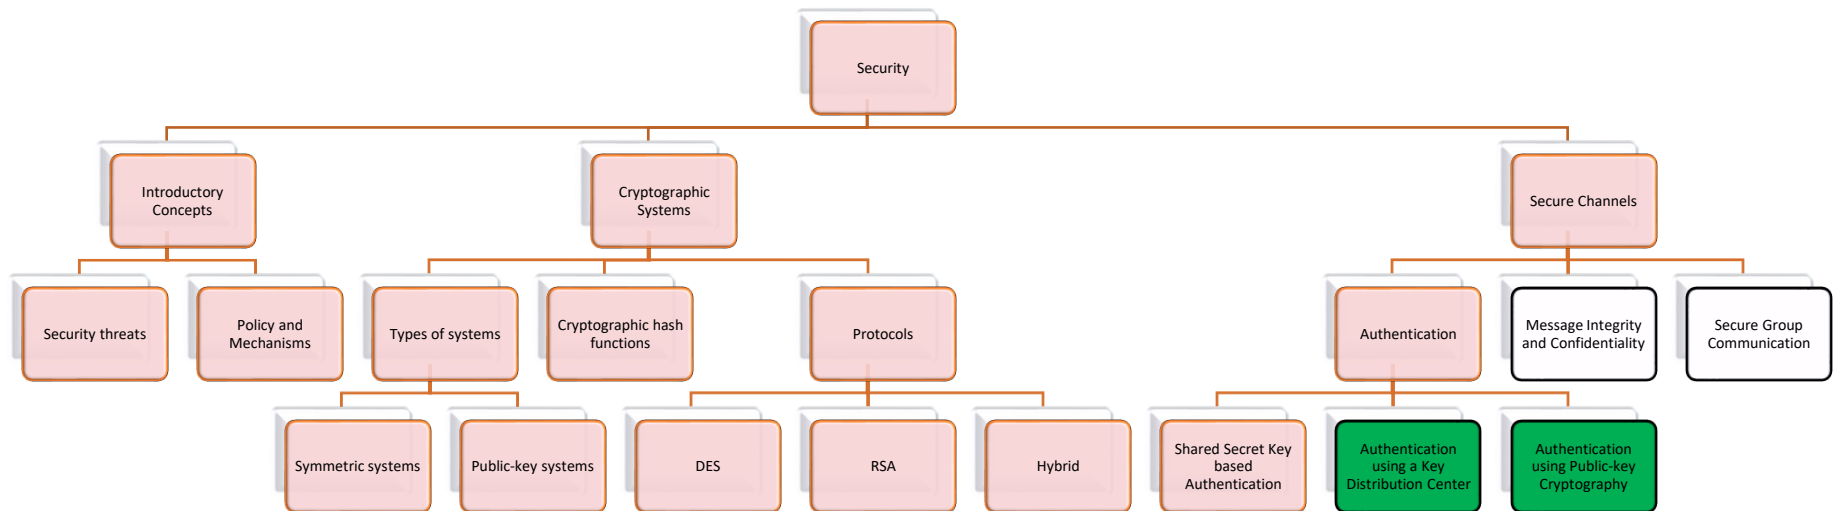
Access Control

- Access Control Matrix
- Protection Domains

Security Management

- Key Management
- Authorization Management

Overview



Types of Mutual Authentication Protocols

1. Shared Secret Key based Authentication
2. Authentication using a Key Distribution Center
3. Authentication using Public-key Cryptography

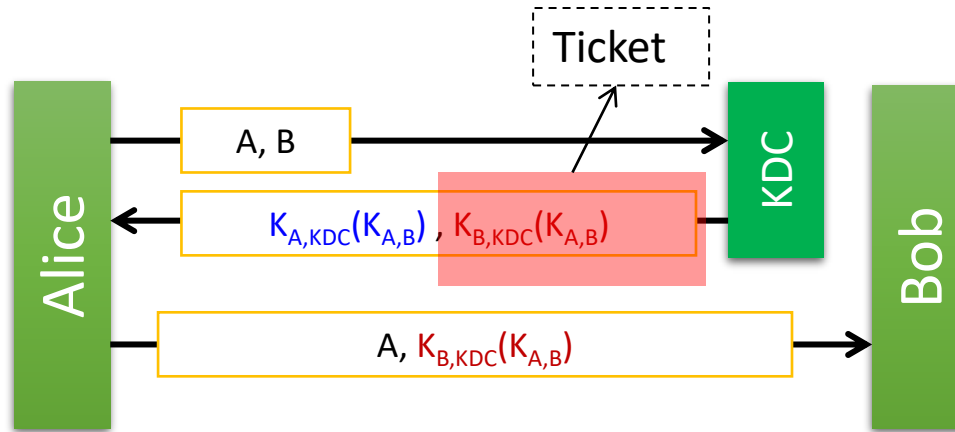
Types of Mutual Authentication Protocols

1. Shared Secret Key based Authentication
2. Authentication Using a Key Distribution Center
3. Authentication Using Public-key Cryptography

Authentication Using a Key Distribution Center

- Shared secret key based authentication is not scalable
 - Each host has exchange keys with every other host
 - Complexity: $O(N^2)$
- The complexity is reduced by electing one node as “*Key Distribution Center*” (KDC)
 - KDC shares a secret-key with every host
 - No pair of hosts need to share a key
 - Complexity: $O(N)$

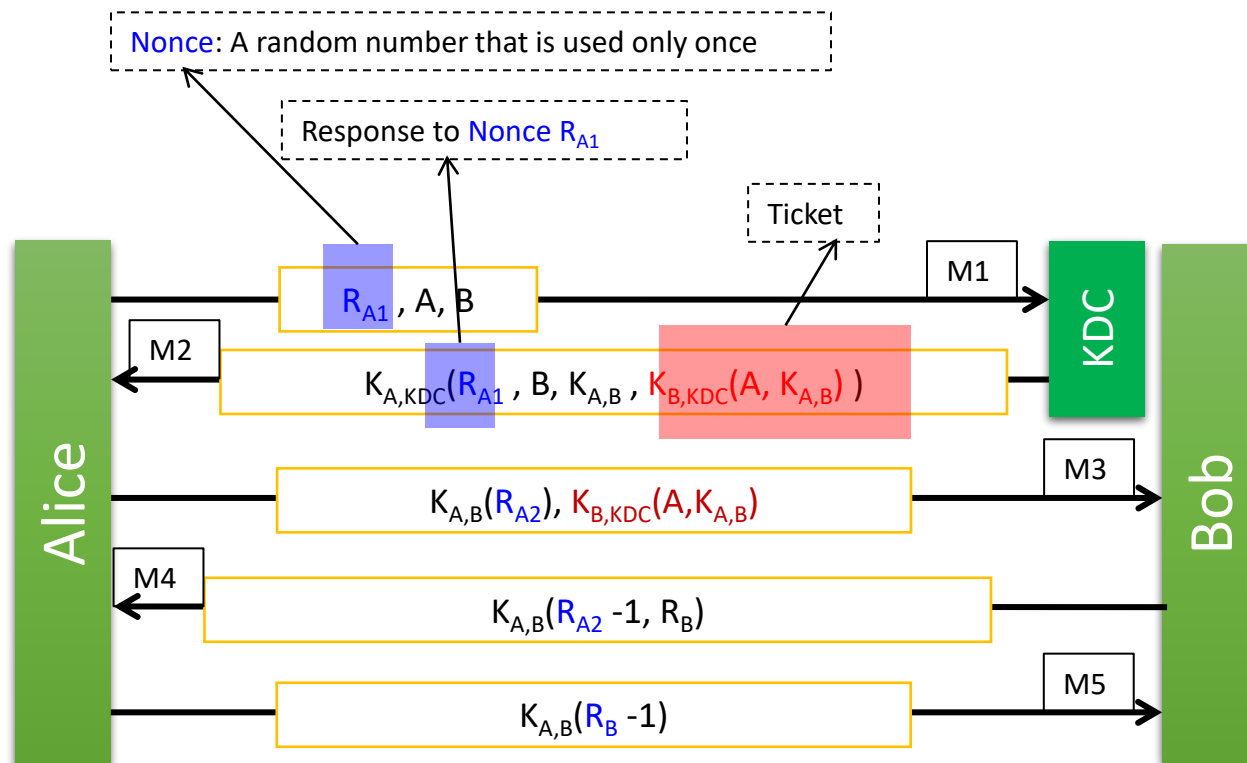
Basic Model of a KDC based Authentication



- We will study a widely used KDC protocol called **Needham-Schroeder Authentication Protocol**

Needham-Schroeder Authentication Protocol

- A multi-way challenge response protocol



What happens if:

1. Nonce is not included in M2?
2. B is not included in M2?
3. R_{A2} is returned in M4, instead of $(R_{A2}-1)$?
4. Chuck has an old key $K_{A,B}$ and intercepts M3, and replays M3 at a later point of time to Alice?

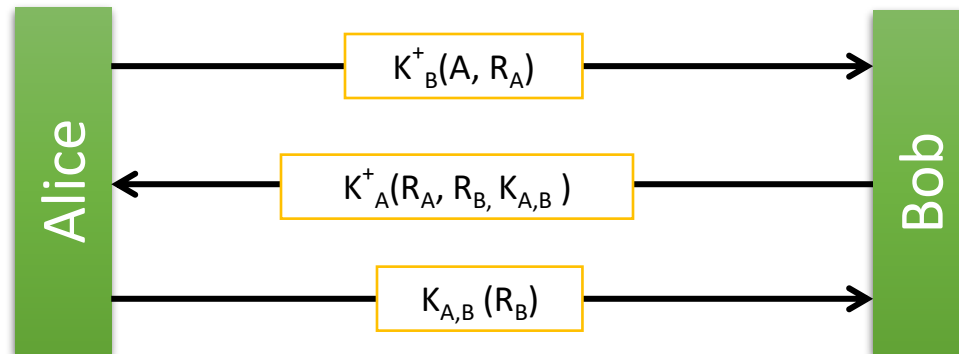
Types of Mutual Authentication Protocols

1. Shared Secret Key based Authentication
2. Authentication using a Key Distribution Center
3. Authentication using Public-key Cryptography

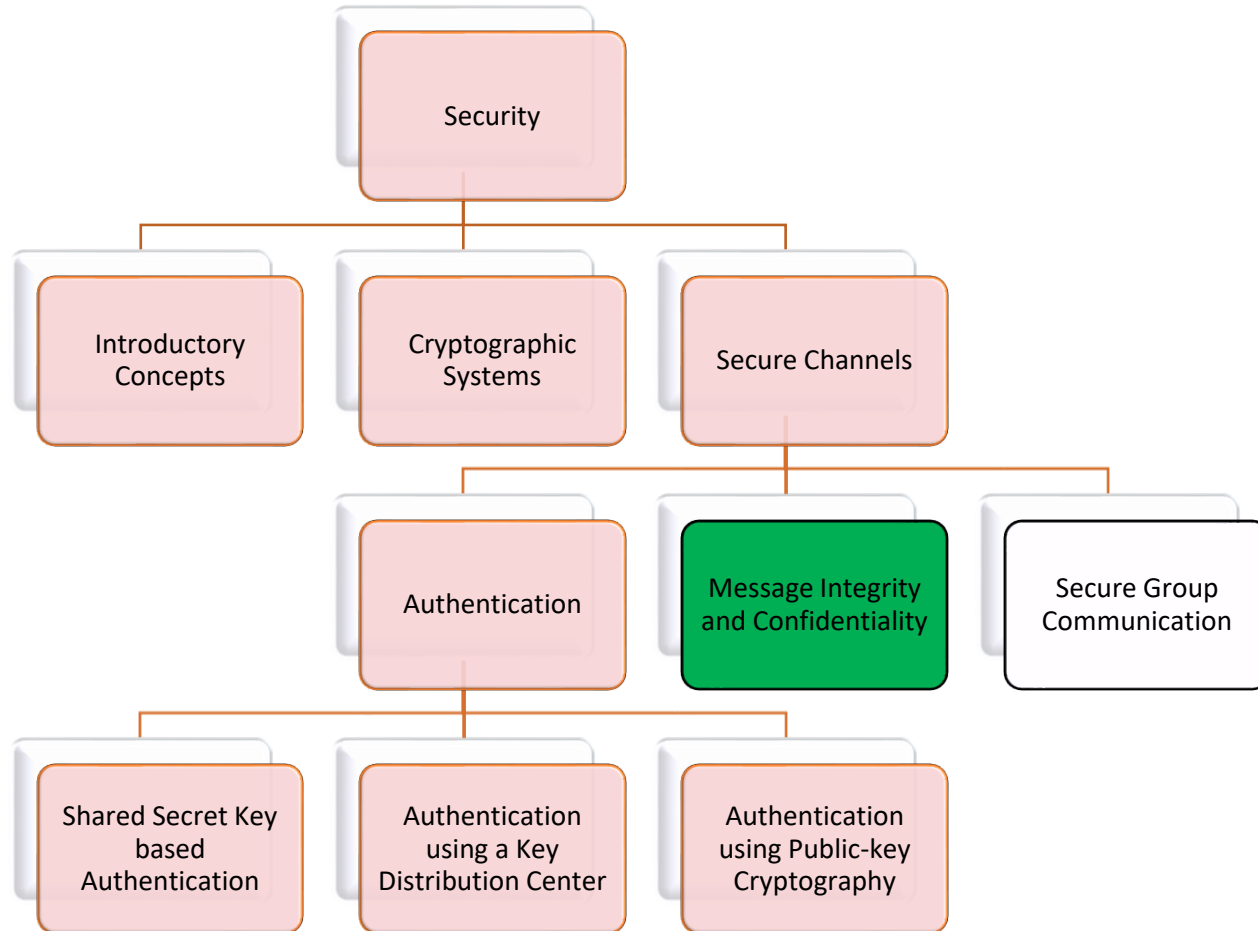
Authentication Using a Public-Key Cryptography

- Recall:

- K_N^+ : Public key of user N
- K_N^- : Private key of user N

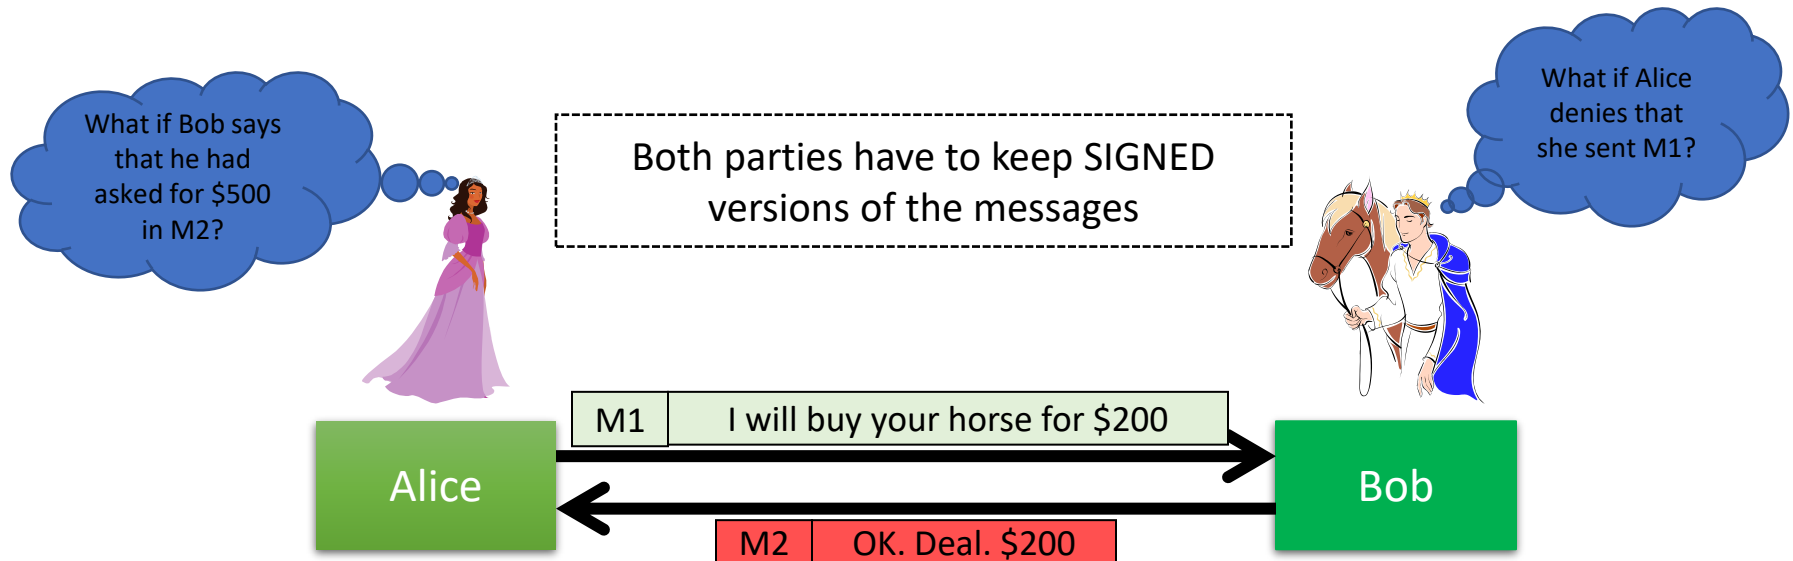


Overview



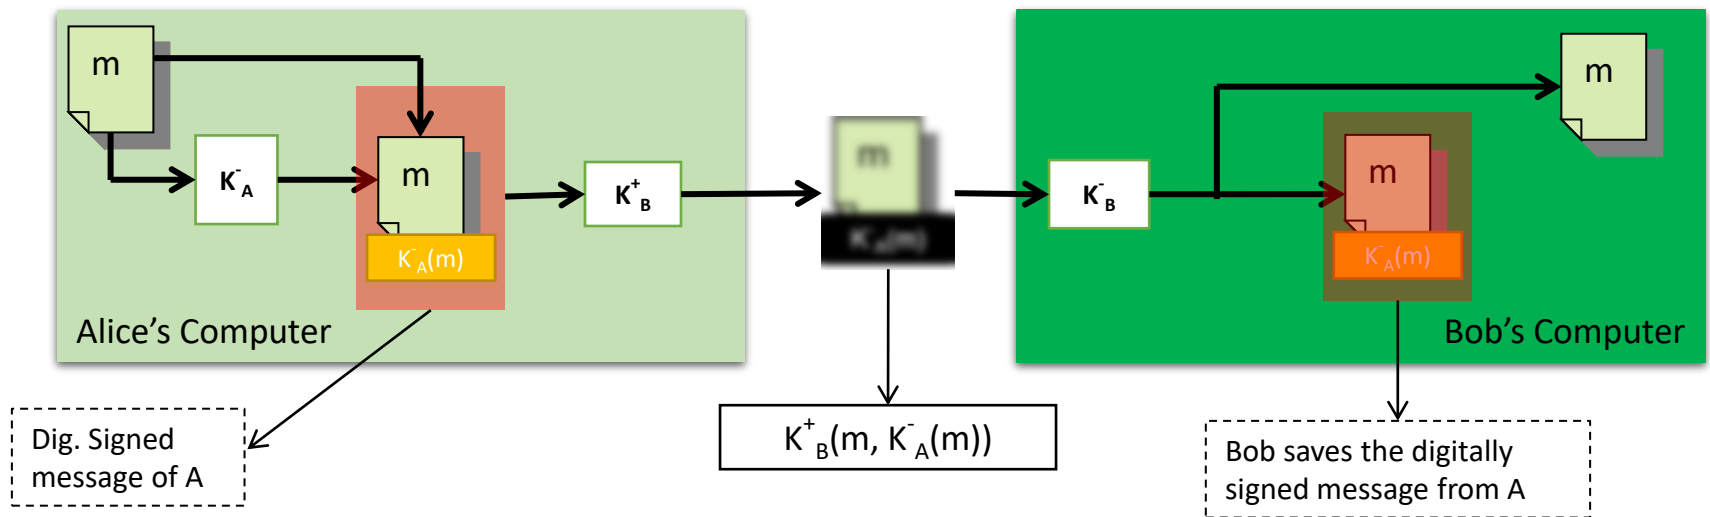
Message Integrity and Confidentiality

- Encryption is used for providing Confidentiality
- But, how to provide Message Integrity?
 - Encryption protects message modification by third party adversaries
 - How to protect modification at senders or receivers?



Digital Signatures

- The sender and the receiver will keep digitally signed copies of the messages
- One way of creating digital signatures is by using RSA

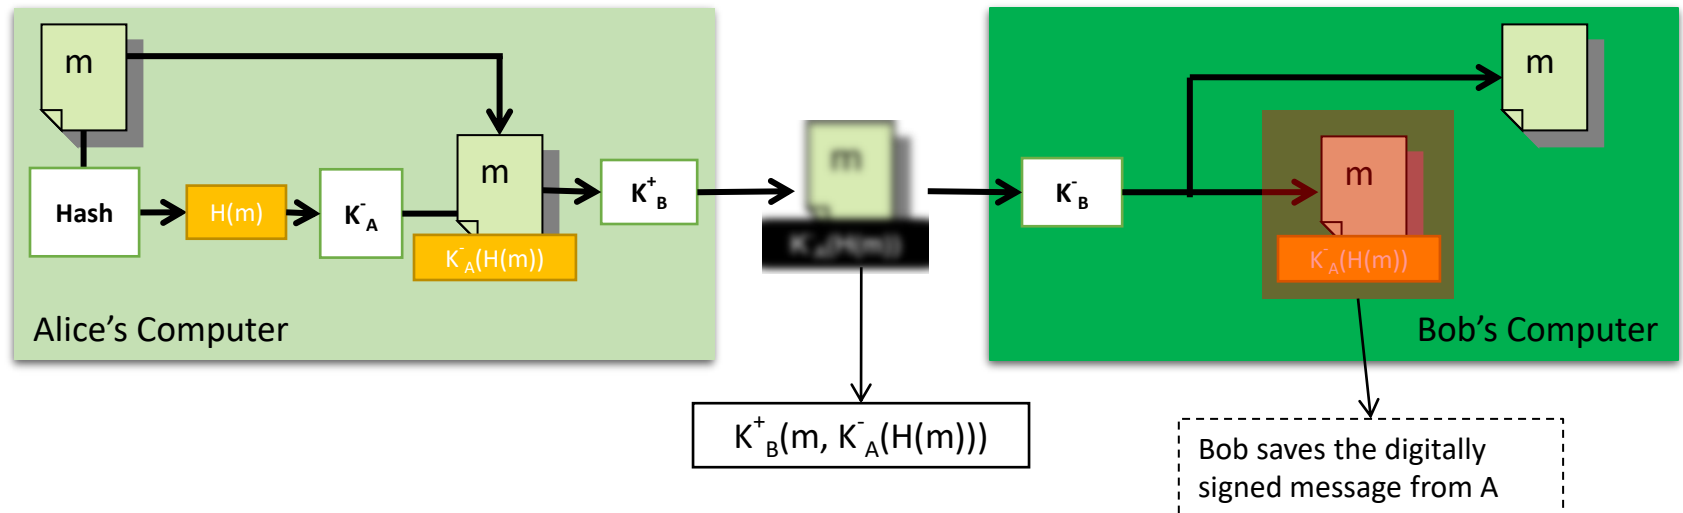


Issues with Naïve Digital Signatures

- The validity of the digitally signed message holds as long as A's private key remains a secret
 - 'A' can claim that her private key was stolen
- If 'A' changes her private key, the signed message becomes invalid
 - A centralized authority can keep track of when keys are changed
- Encrypting the entire message with private key is costly
 - RSA is slower, and encryption of long data is not preferred
 - Any ideas on how to improve?

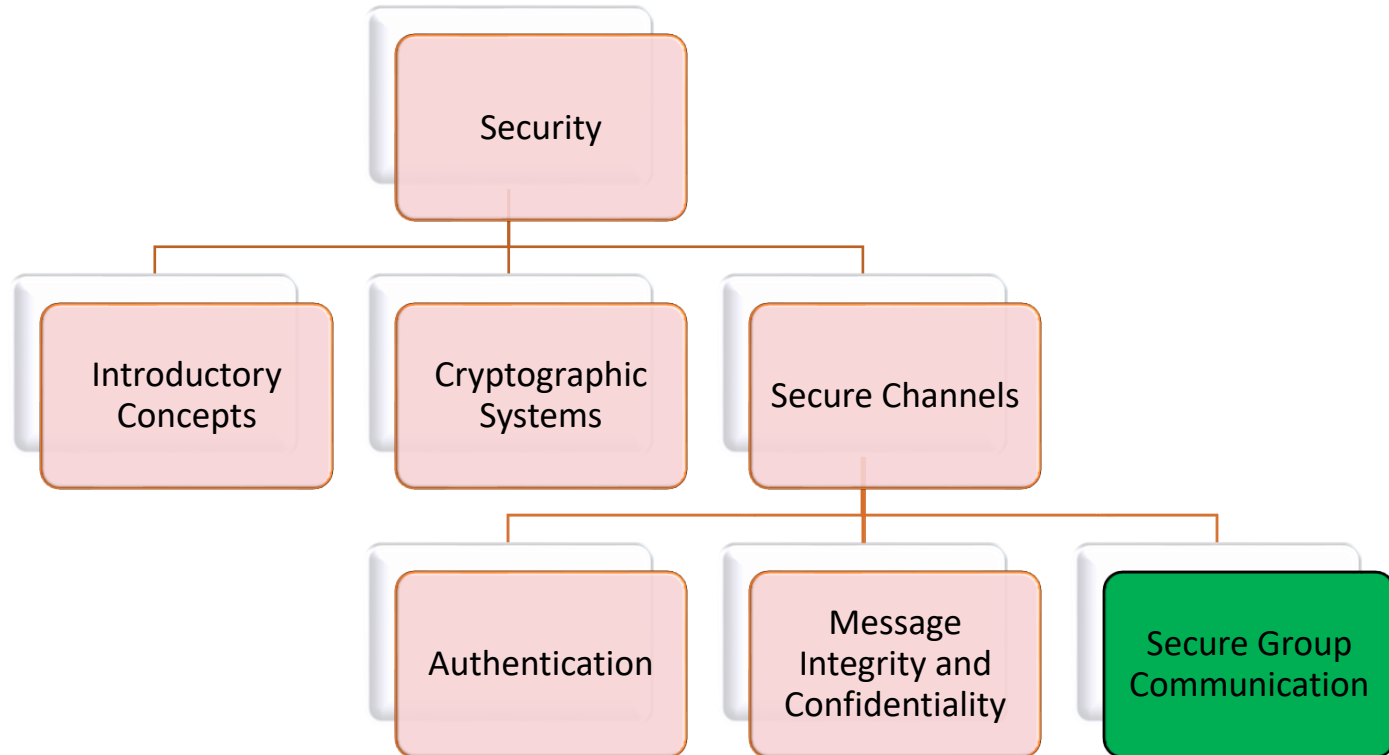
Digital Signature with Message Digest

- Instead of attaching the encrypted message, attach the *message digest*



Bob can verify the signature by comparing the hash of received message, and the hash attached in the digitally signed message

Overview



Secure Group Communication

- Scenario: Client has to communicate with a group of replicated servers
 - Replication can be for fault-tolerance or performance improvement
- Client expects that the response is secure
- Naïve Solution:
 - Client collects the response from each server
 - Client authenticates every response
 - Client takes a majority vote
- Drawback of naïve approach
 - Replicated of servers should be transparent to the client
 - Reiter et al. proposed a transparent and secure group authentication protocol ([Secure Group Authentication Protocol](#))

Secure Group Authentication Protocol

- Problem: Authenticate a group of replicated servers at the client in a transparent way
- Main Idea:
 - Multiple servers share a secret
 - None of them know the entire secret
 - Secret can be revealed only if all/most of them cooperate
- The approach is similar to k-fault tolerant systems
 - Tolerate intrusion (faults) of c servers out of N servers if $c \leq N$
- We will study the algorithm in two phases:
 1. Generating Shared Secret Signatures
 2. Enhancing transparency

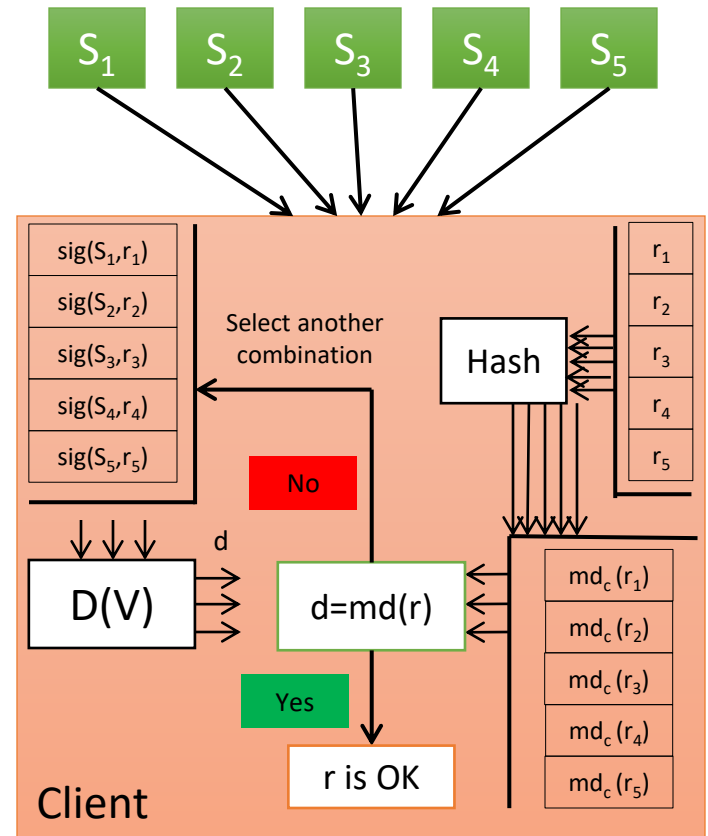
1. Generating Shared Secret Signatures

- For each client call:
 - Each server S_i returns the response (r_i)
 - In addition, it sends a signed message digest: $\text{sig}(S_i, r_i) = K_{S_i}^- (\text{md}(r_i))$
- For each response, client computes $\text{md}_c(r_i) = H(r_i)$
- Hence, client has N triplets $R_i = \langle r_i, \text{md}_c(r_i), \text{sig}(S_i, r_i) \rangle$
- Client takes $c+1$ combinations of triplets
- For each combination vector $[R, R' \text{ and } R'']$
 - Computes a create a single digest using a special decryption function $D(V)$

$$d = D(V) = D(\text{sig}(S, r), \text{sig}(S', r'), \text{sig}(S'', r''))$$

NOTE: d is a vector of responses
 - If $d[x] = \text{respective } \text{md}_c(r_i)$, then r_i is correct

Parameters: $N=5; c=2$



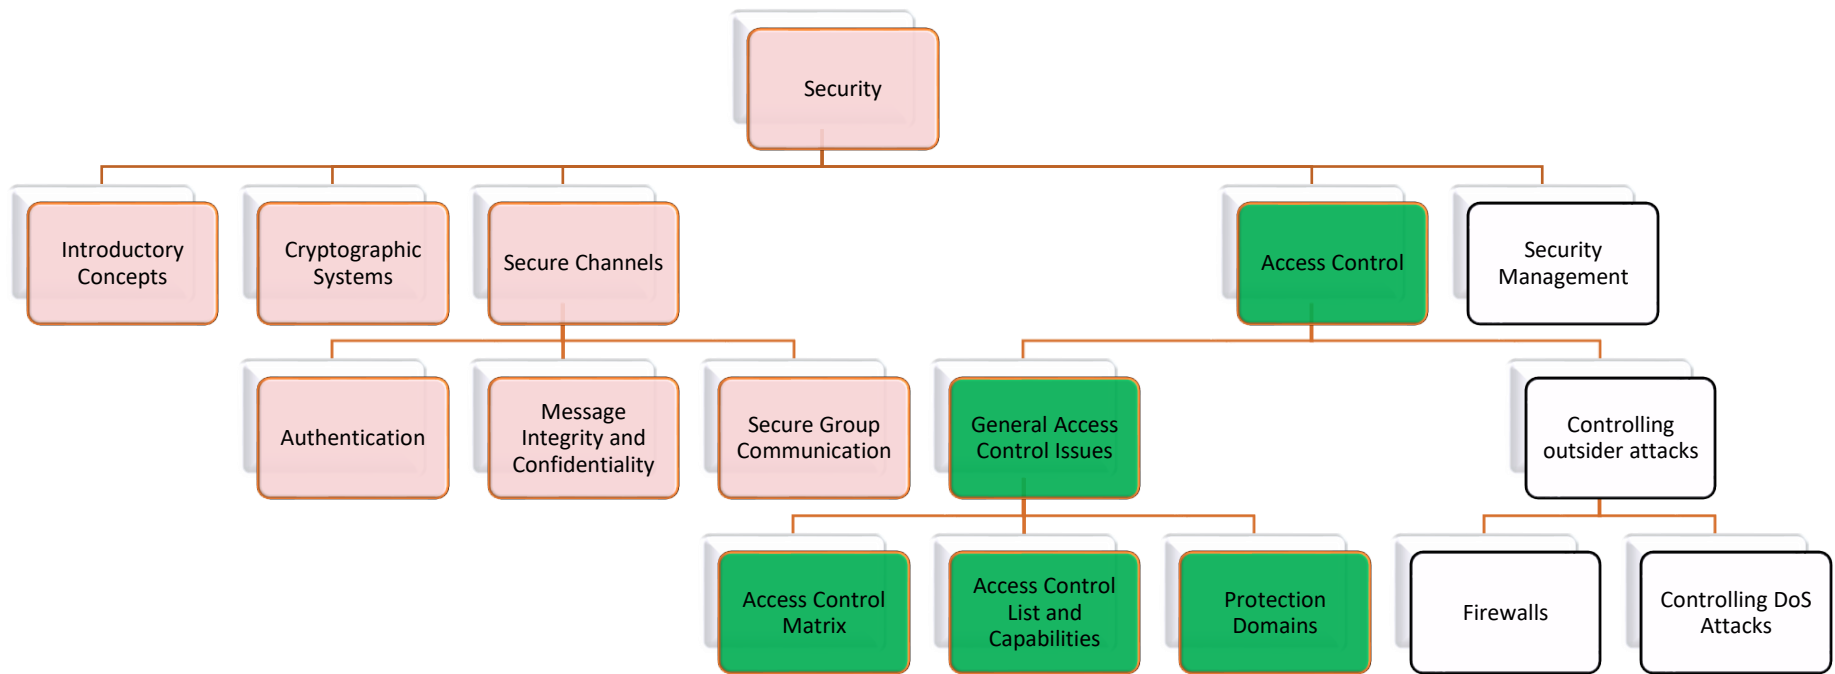
2. Enhancing transparency

- Each server broadcasts a message $[r_i, \text{sig}(S_i, r_i)]$ to other servers
- When the server has received $k+1$ messages
 - Compute valid signature $D(V)$ for one combination of responses r
 - If success, then transmit $[V, r]$ to the client
- When client receives a response, it re-verifies the correctness of r

Discussion: Secure Group Authentication Protocol


- The protocol can be generalized to (m,n) -threshold scheme
 - A message is divided into n -pieces (shadows)
 - m shadows can be used to reconstruct the message
 - $(m-1)$ or fewer shadows cannot reconstruct
- In the above protocol:
 - $m = c+1$ = Number of valid signatures required
 - $n = N$ = Number of replicated servers

Overview




Access Control

- Access Control (AC) verifies the access rights of a subject requesting an access to an object




- A General model for Access Control



- A reference monitor records which subjects may do what
- How is Access Control different from Authorization?
 - Authorization = granting access rights
 - Access Control = verifying access rights


Access Control Matrix

- A matrix M that keeps track of which operations can a subject invoke on an object
 - $M[s,o]$ lists which operations subject s can perform on object o
- The reference monitor will look up AC Matrix before authorizing the request
- Scalability of AC Matrix
 - For a large DS with many users, most of the entries are empty
- Two implementations of AC Matrix:
 1. Access Control Lists
 2. Capabilities




1. Access Control List

- AC Matrix is stored as a list
 - Similar to how a graph of nodes and edges are stored as an adjacency list



2. Capabilities

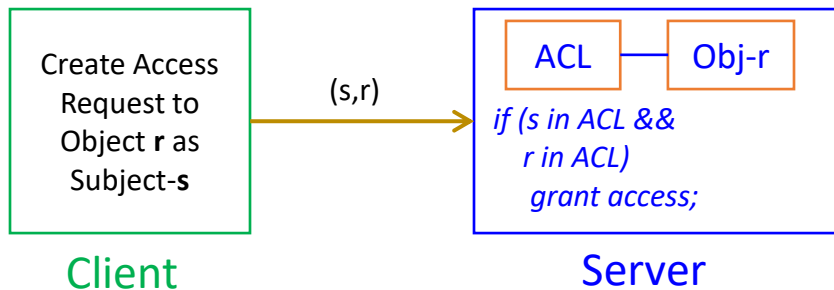
- AC Matrix can also be decomposed row-wise
- Capability for a user is a list of objects and the associated permissions for a given user



- The 'capability' object can be issued by Reference Monitor to a user, and can reside at the user side
 - But should be protected (digital signature)


Comparison of ACL and Capabilities

Using ACL



The server stores the ACL

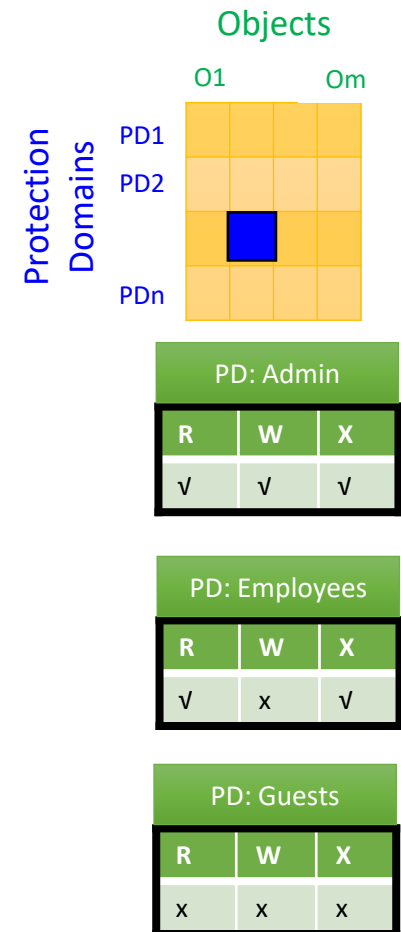
Using Capabilities



- The server is not interested in whether it knows the client
- Server should check Capability

Protection Domains

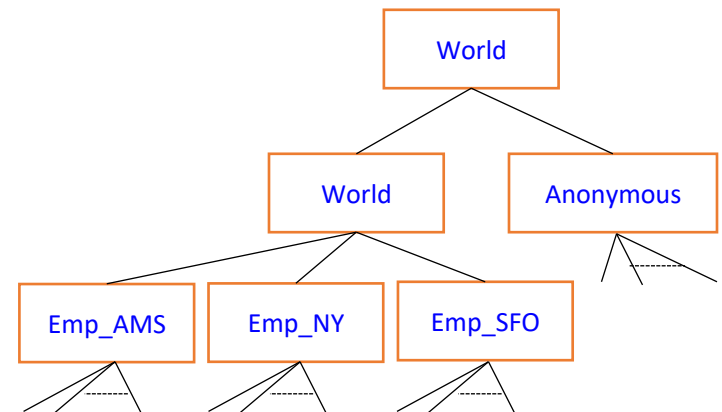
- Motivation: Access Control using ACL and Capabilities is still not scalable in very large DS
- Organizing Permissions:
 - Object permission is given to a protection domain (PD)
 - Each PD row consists of (object, permission) pair
- User Management:
 - Group users into a “protection domain”
 - Users belong to one or more protection domain
 - e.g., employees, admin, guests
 - Before granting the access, check permissions from user’s PDs



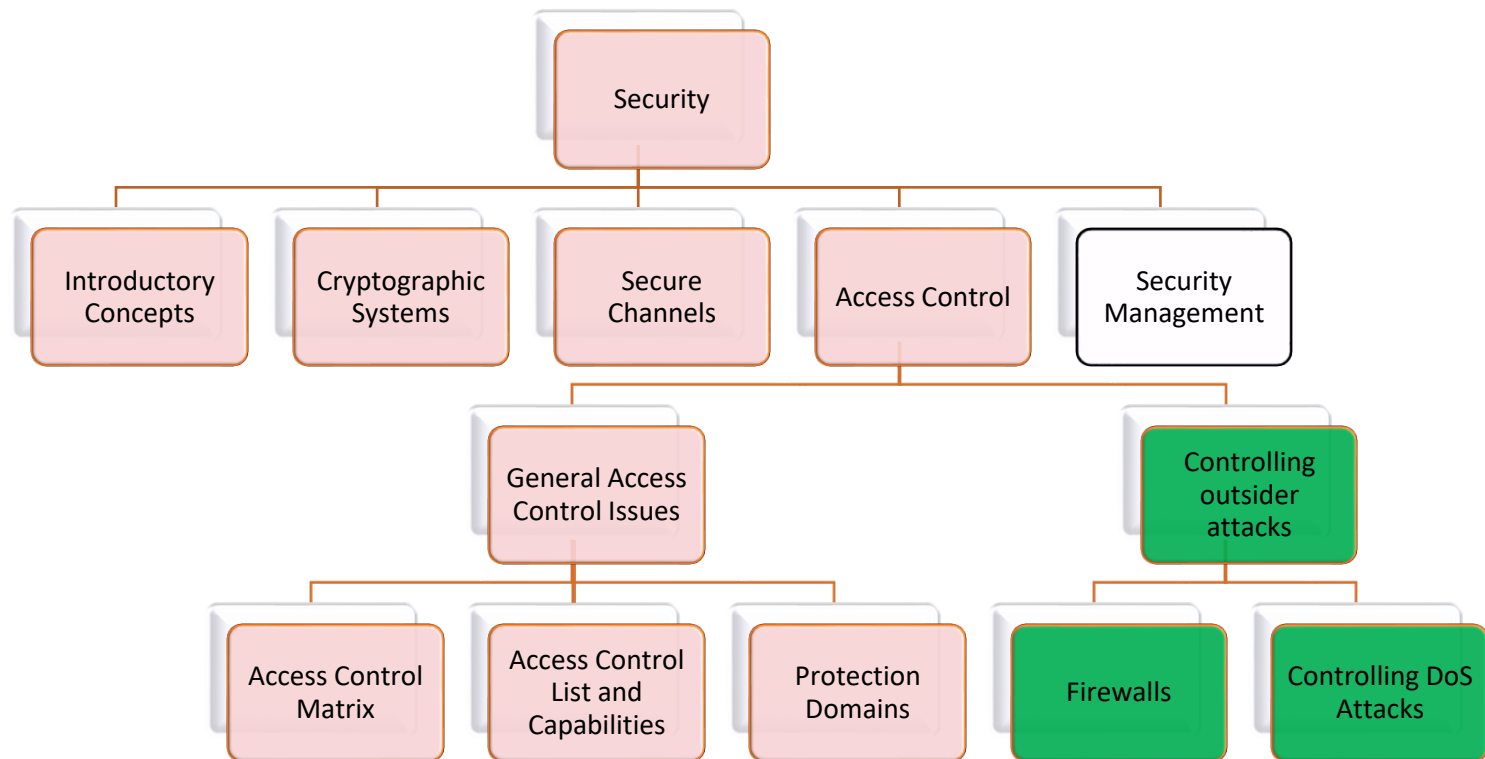
A matrix is shown only for representation. Similar to AC Matrix, PD can be implemented using techniques similar to ACL

Hierarchical Protection Domains

- Domains can be organized into hierarchical groups
- Advantage
 - Managing group membership is easy
- Disadvantage
 - Looking up for a member is costly
- An optimization:
 - Let each user carry a certificate of all the groups they belong to
 - Reference monitor validates the certificate (similar to the Capabilities approach)



Overview




Controlling Outsider Attacks

- Above discussed schemes are good for preventing against small number of users who belong to a DS
- But, what if any user on the Internet can access the resources of DS
 - Search engines, Sending mails, ...
- Above approaches are complicated, inefficient and time-consuming when a DS is exposed to the rest-of-the-world
- We discuss two coarse grained Access Control schemes
 - Firewall
 - Controlling Denial of Service (DoS) attacks

Firewall

- Firewall provides access control by regulating the flow of network traffic to/from a LAN
- Firewall disconnects any part of the DS from outside world
 - The internal LANs route packets through Packet Filtering Routers (PFR) and Application Gateway (AG)
 - Incoming/outgoing packet contents (e.g., src/dest IP) are examined by PFRs
 - PFRs selectively forward the incoming/outgoing packets to/from AG



Denial of Service Attacks

- *Denial of Service* (DoS) attacks maliciously prevents authorized processes from accessing the resources
 - Example: Flood a search server with TCP requests
- DoS attacks from a single source can be easily prevented
- Triggering DoS attacks from multiple sources (Distributed DoS or DDoS) is hard to detect
 - Attackers hijack a large group of machines
 - DDoS is hard to detect

Types of DoS Attacks

Two types

1. Bandwidth Depletion

- Send many messages to a server
- Normal messages cannot reach the server

2. Resource Depletion

- Trick the server to allocate large amount of resources
- Example: TCP SYN flooding
 - Initiate a lot of TCP connections on server (by sending TCP SYN packets)
 - Do not close the connections
 - Each connection request blocks some memory for some time

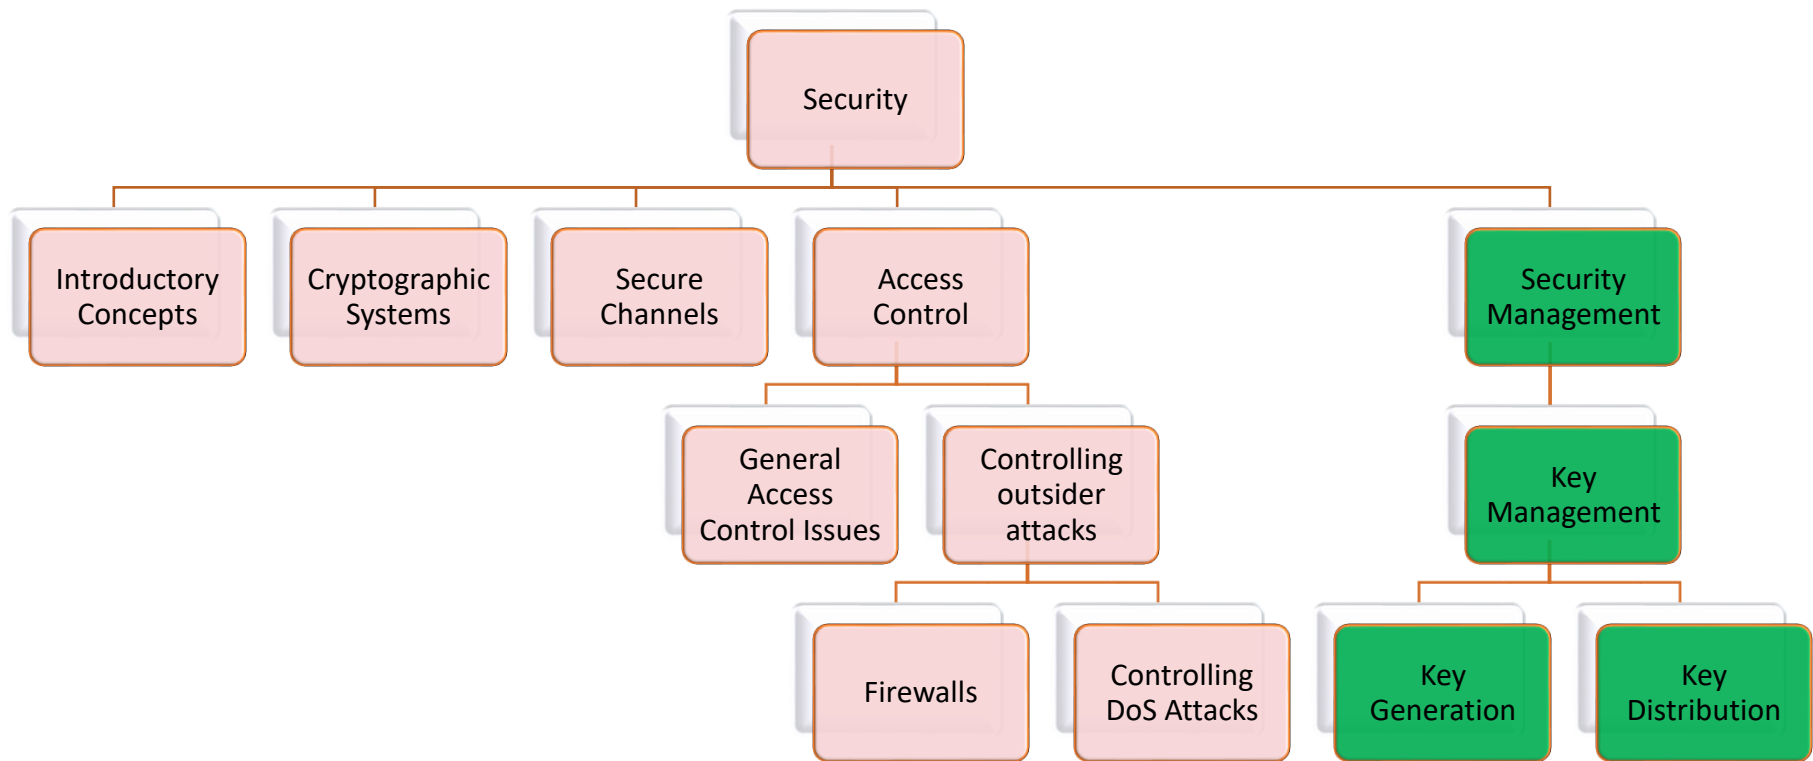
Preventing DoS Attacks

Idea: Detect DDoS by continuously monitoring network traffic at routers

Three types:

- Monitor at ingress routers:
 - Monitor incoming traffic for large flow towards a destination
 - Drawback: Too late to detect since regular traffic is already blocked
- Monitor at egress routers:
 - Monitor traffic when packets leave organization's network
 - e.g., Drop packets whose source IP does not belong to the org's network
- Monitor at core internet routers:
 - Monitor at routers at the core of the Internet Service Providers (ISPs)
 - e.g., Drop packets when rate of traffic to a specific node is disproportionately larger than the rate of traffic from that node

Overview




Key Management


- We have studied how to use keys for authentication and encryption
- Examples of key-generation that we studied:
 - Session key
 - Generated and distributed by public-private key pair
 - Public-Private Key
 - Generated by algorithms. How is it securely distributed?
- How to boot-strap secure key generation and distribution in a DS?
- There are two main approaches in key management:
 - Diffie-Hellman Algorithm (For key generation over insecure channels)
 - Key distribution

Diffie-Hellman Key Generation

- Diffie-Hellman algorithm establishes a shared key between two entities across an insecure channel



Diffie-Hellman Algorithm

Alice	Eve	Bob
Selects two large numbers n and g with certain mathematical properties		
Choose a secret number x		Choose a secret number y
Alice transmits $[n, g, g^x \bmod n]$	$n, g, g^x \bmod n$	$y, n, g, g^x \bmod n$
$n, g, x, g^x \bmod n, g^y \bmod n$	$n, g, g^x \bmod n, g^y \bmod n$	Bob transmits $[g^y \bmod n]$
Alice computes $(g^y \bmod n)^x = g^{xy} \bmod n$		Bob computes $(g^x \bmod n)^y = g^{xy} \bmod n$
SHARED SECRET KEY = $g^{xy} \bmod n$		SHARED SECRET KEY = $g^{xy} \bmod n$

Discussion of Diffie-Hellman Algorithm

- D-H algorithm can be used in:
 - Symmetric Cryptosystems
 - Secret key $g^{xy} \bmod n$ can be generated between A and B even in the presence of insecure channel
 - Public-key Cryptosystems
 - A's private key = x
 - A's public key = $g^x \bmod n$
- But, how to securely distribute public keys?


Key Distribution

- Key distribution in Symmetric Cryptosystems
 - Out-of-band key exchange
 - Diffie-Hellman algorithm
 - Complex for a large DS
- Key distribution in Public-key Cryptosystems
 - Main problem: How do we authenticate the entity that transmitted the public key?
 - We will study an approach called Public-key certificates

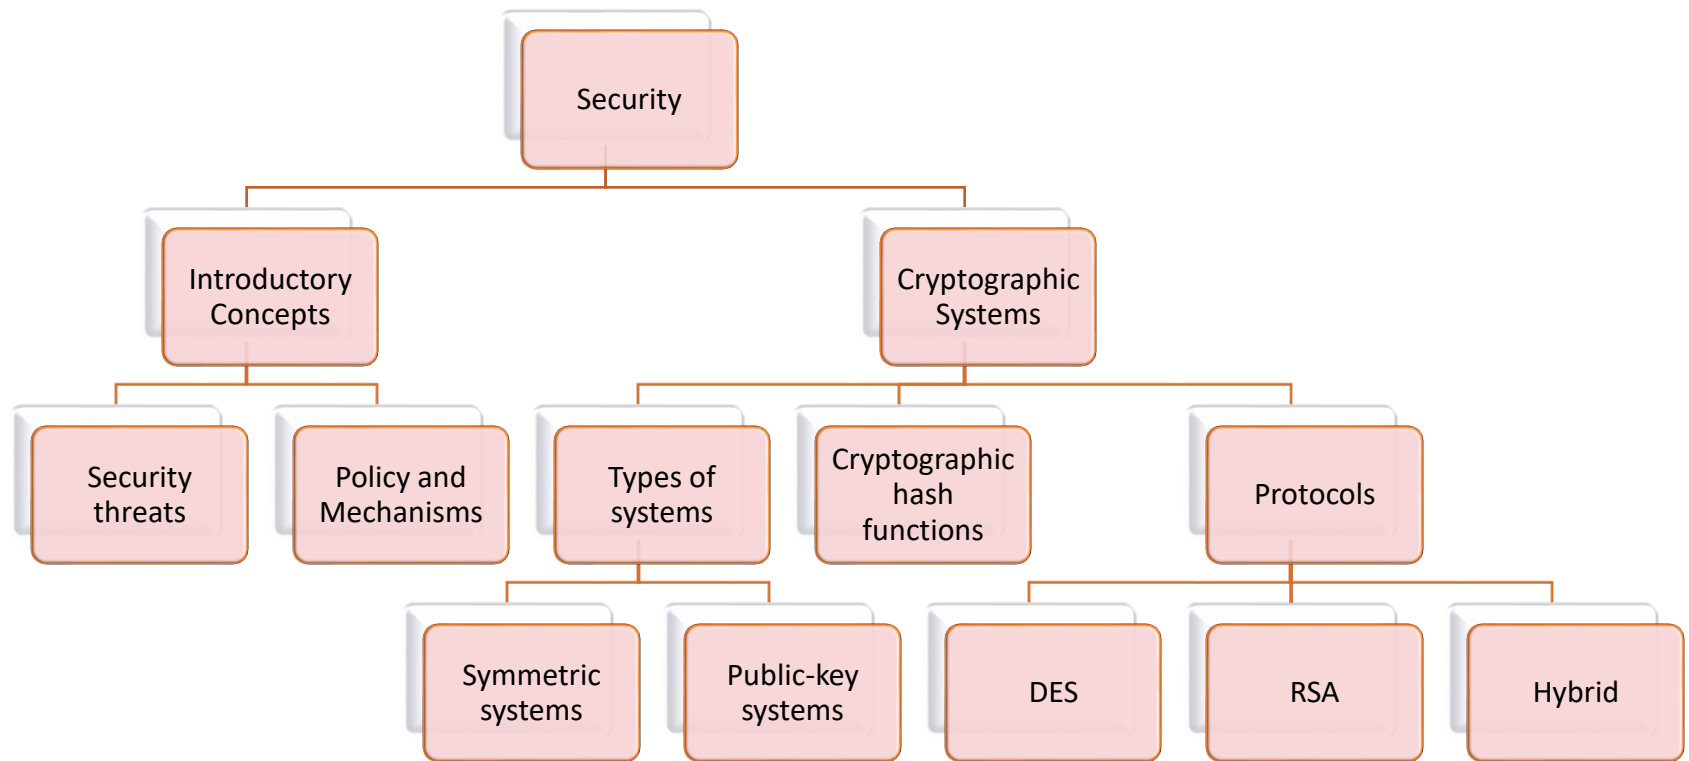
Public-key Certificate

- In practice, we trust a well-know Certification Authority (CA)
 - Example of a CA: VeriSign
- Public key of CA (K_{CA}^+) is well-known
 - e.g., It is built-in with browsers
- How does Bob validate public key of Alice?
 - Each user (Alice) has a public-key certificate $\text{cert}(A, K_A^+)$ that issued by CA
 - Bob verifies if the received (A, K_A^+) matches that in the certificate

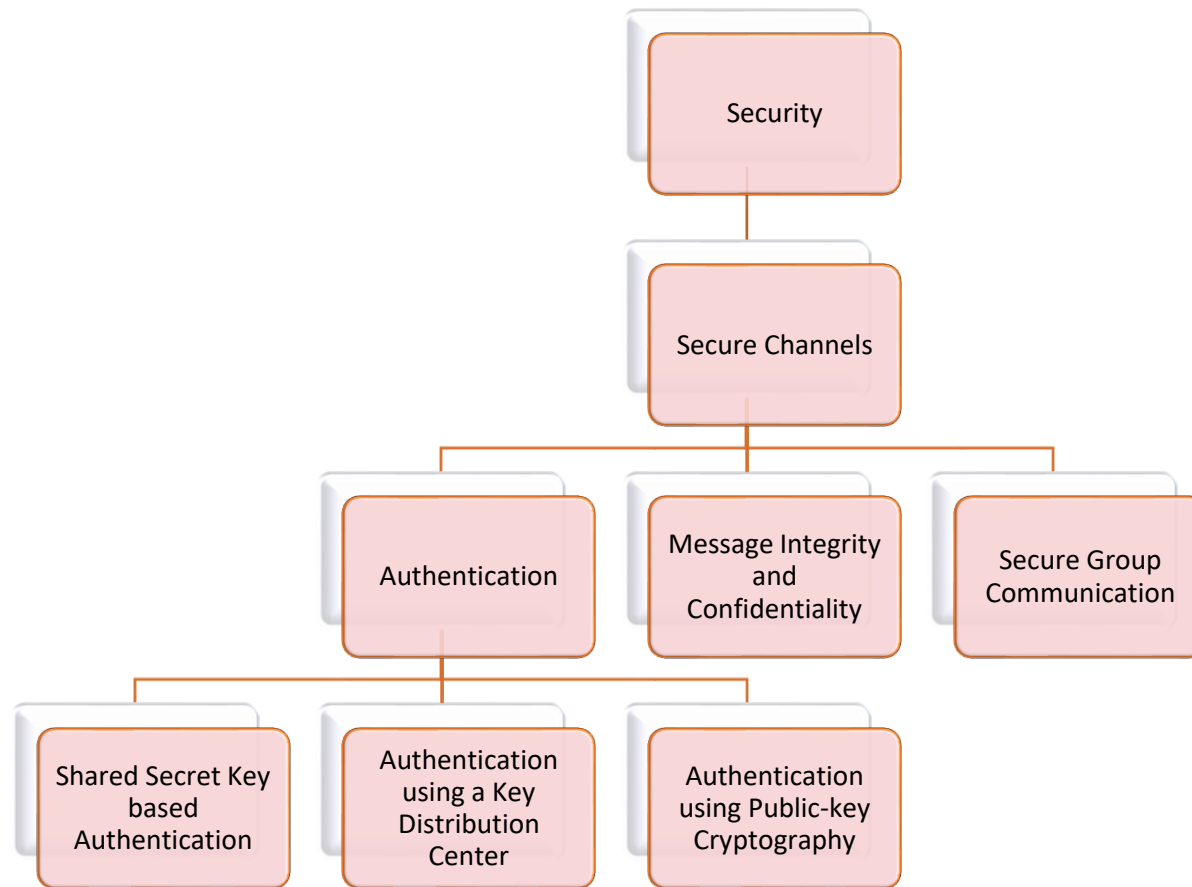
High-level Summary of Security Chapter



Summary of “Introductory Concepts” and “Cryptographic Systems”



Summary of “Secure Channels”



Summary of Security Chapter

